

CSE463: Neural Networks

“Computer Vision and Image Processing: Recent Advances and Trends”

Deep Learning

by:

Hossam El Din Hassan Abd El Munim

حسام الدين حسن عبد المنعم

Computer & Systems Engineering Dept.,

Ain Shams University,

1 El-Sarayat Street, Abbassia, Cairo 11517

Essential Books

1. **Ian Goodfellow and Yoshua Bengio and Aaron Courville, Deep Learning, MIT Press, 2016.**
2. **I Charu C. Aggarwal, Neural Networks and Deep Learning, Springer, 2018.**
3. **I Michael Nielsen, Neural Networks and Deep Learning.
<http://neuralnetworksanddeeplearning.com/>, 2016**
4. Rudolf Kruse, C. Borgelt, F. Klawonn, C. Moewes, M. Steinbrecher, P. Held, Computational Intelligence: A Methodological Introduction, Springer, 2014.
(available online)
5. S. Haykin, Neural Networks and Learning Machines, 2009.

Course Contents

- **Basics:** Biological Neuron, Idea of computational units, McCulloch{Pitts unit and Thresholding logic, Linear Perceptron, Perceptron Learning Algorithm, Linear Separability. Convergence theorem for Perceptron Learning Algorithm.
- **Feedforward Networks:** Multilayer Perceptron, Gradient Descent, Backpropagation, Empirical Risk Minimization, regularization, auto-encoders.
- **Deep Neural Networks:** Difficulty of training deep neural networks, Greedy layer-wise training.
- **Better Training of Neural Networks:** Newer optimization methods for neural networks (Adagrad, adadelata, rmsprop, adam, NAG), second order methods for training, Saddle point problem in neural networks, Regularization methods (dropout, drop connect, batch normalization).
- **Recurrent Neural Networks:** Back propagation through time, Long Short Term Memory, Gated Recurrent Units, Bidirectional LSTMs, Bidirectional RNNs
- **Convolutional Neural Networks:** LeNet, AlexNet.
- **Generative models:** Restrictive Boltzmann Machines (RBMs), Introduction to MCMC and Gibbs Sampling, gradient computations in RBMs, Deep Boltzmann Machines.

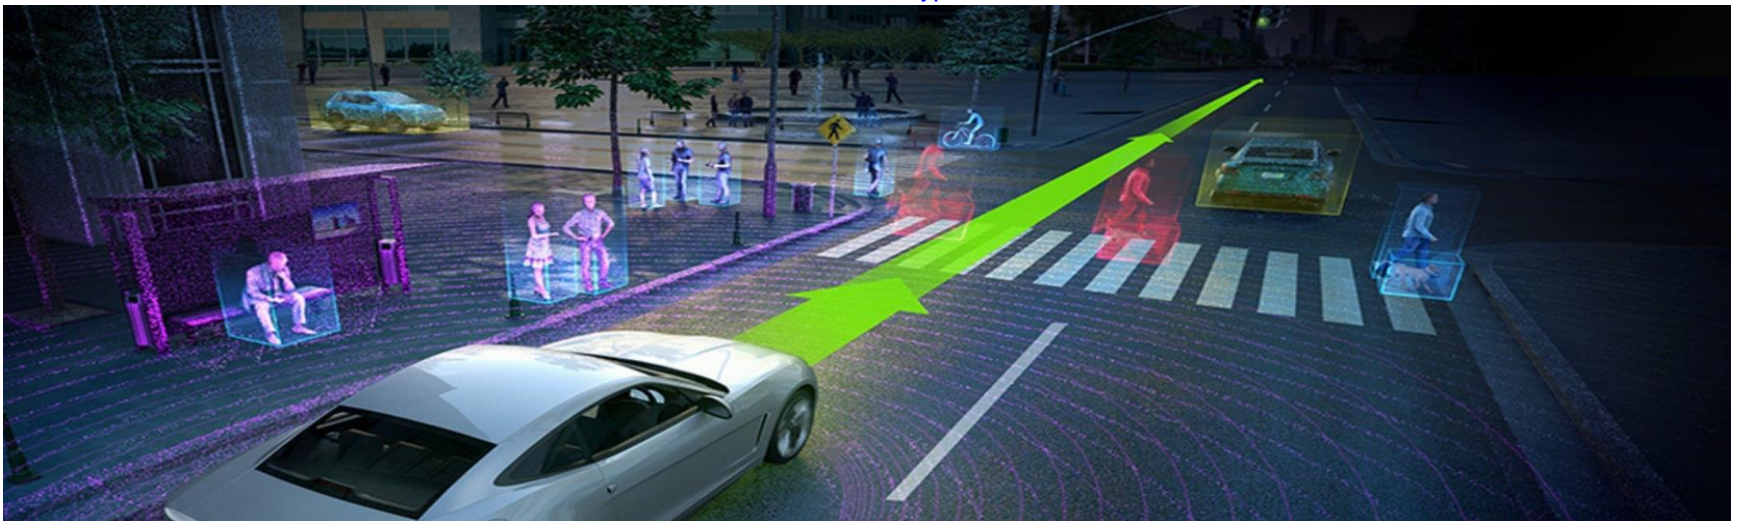
Grading System

- **Final examination 90**
- **Midterm examination 15**
- **Attendance 5**
- **Quizzes (TWO) 5**
- **Project 10**
- **Warnings:**
 - A quiz may be given without being informed before.
 - Copying assignment is prohibited.
 - Delay of submission influences on marks.
 - No Plagiarism!

Autonomous Driving

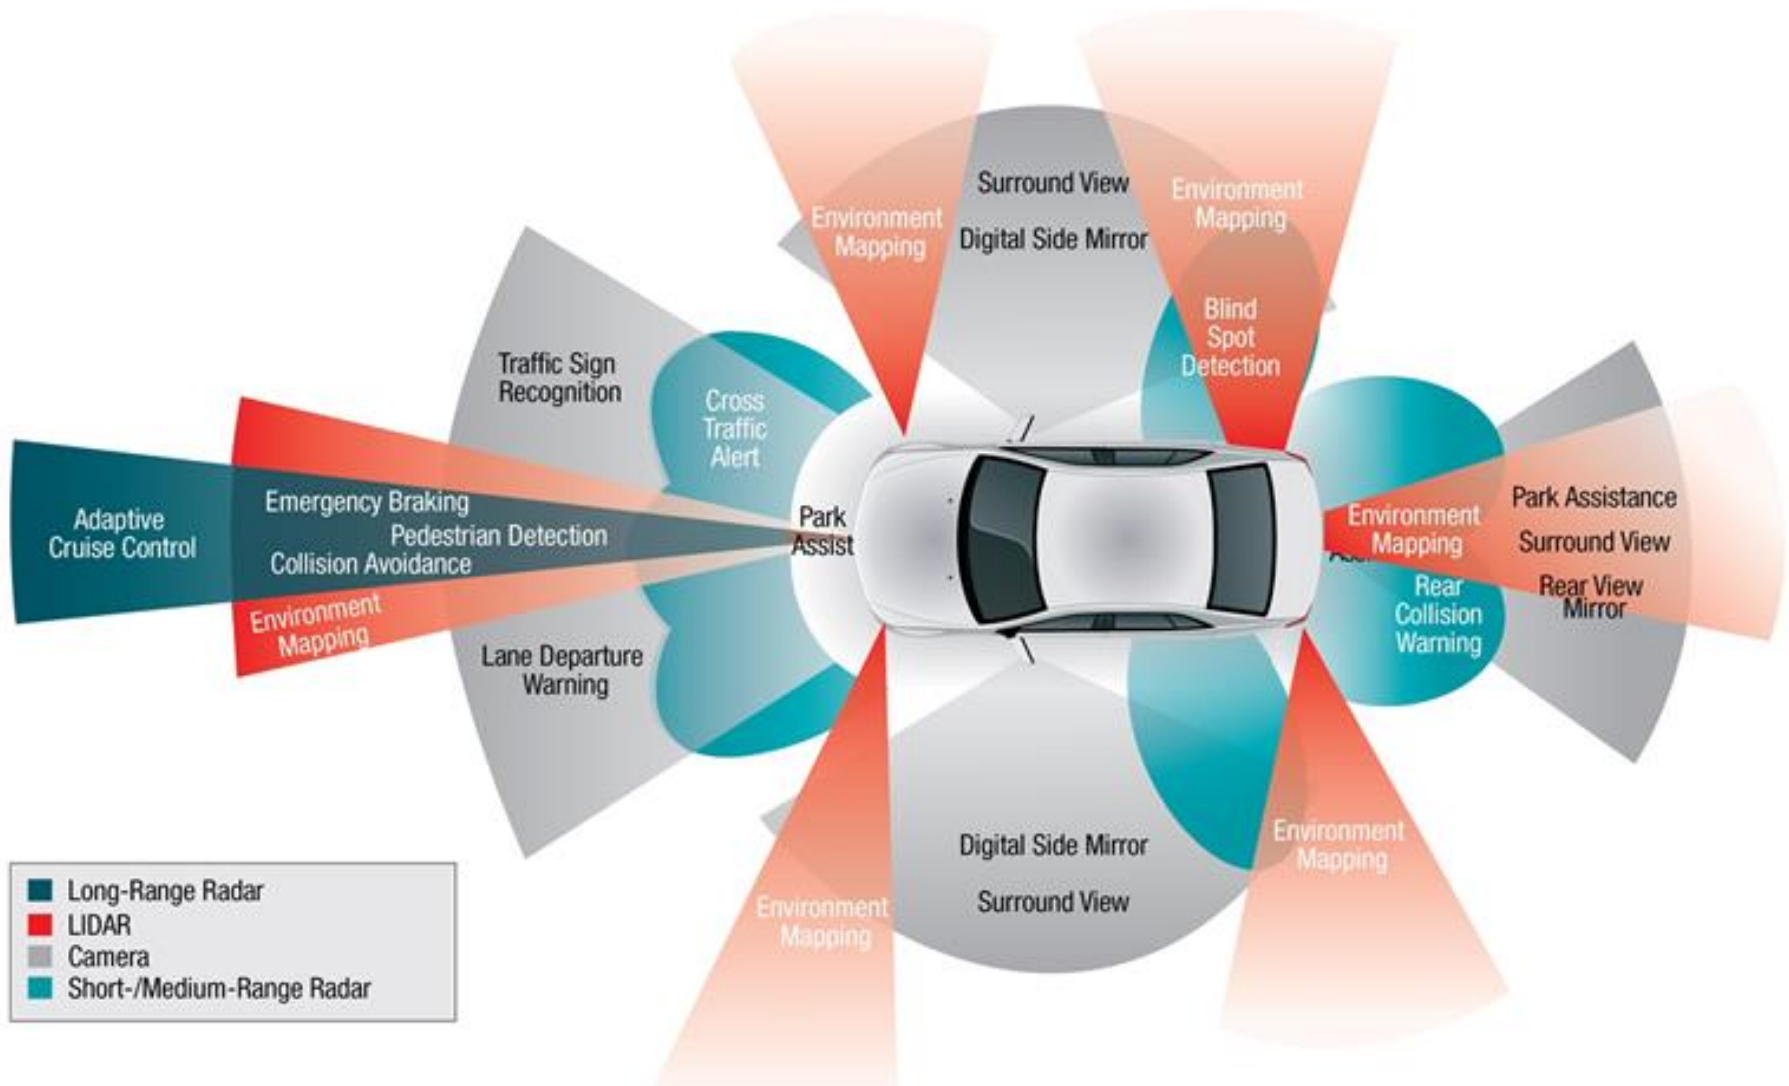


Photo: www.shellypalmer.com



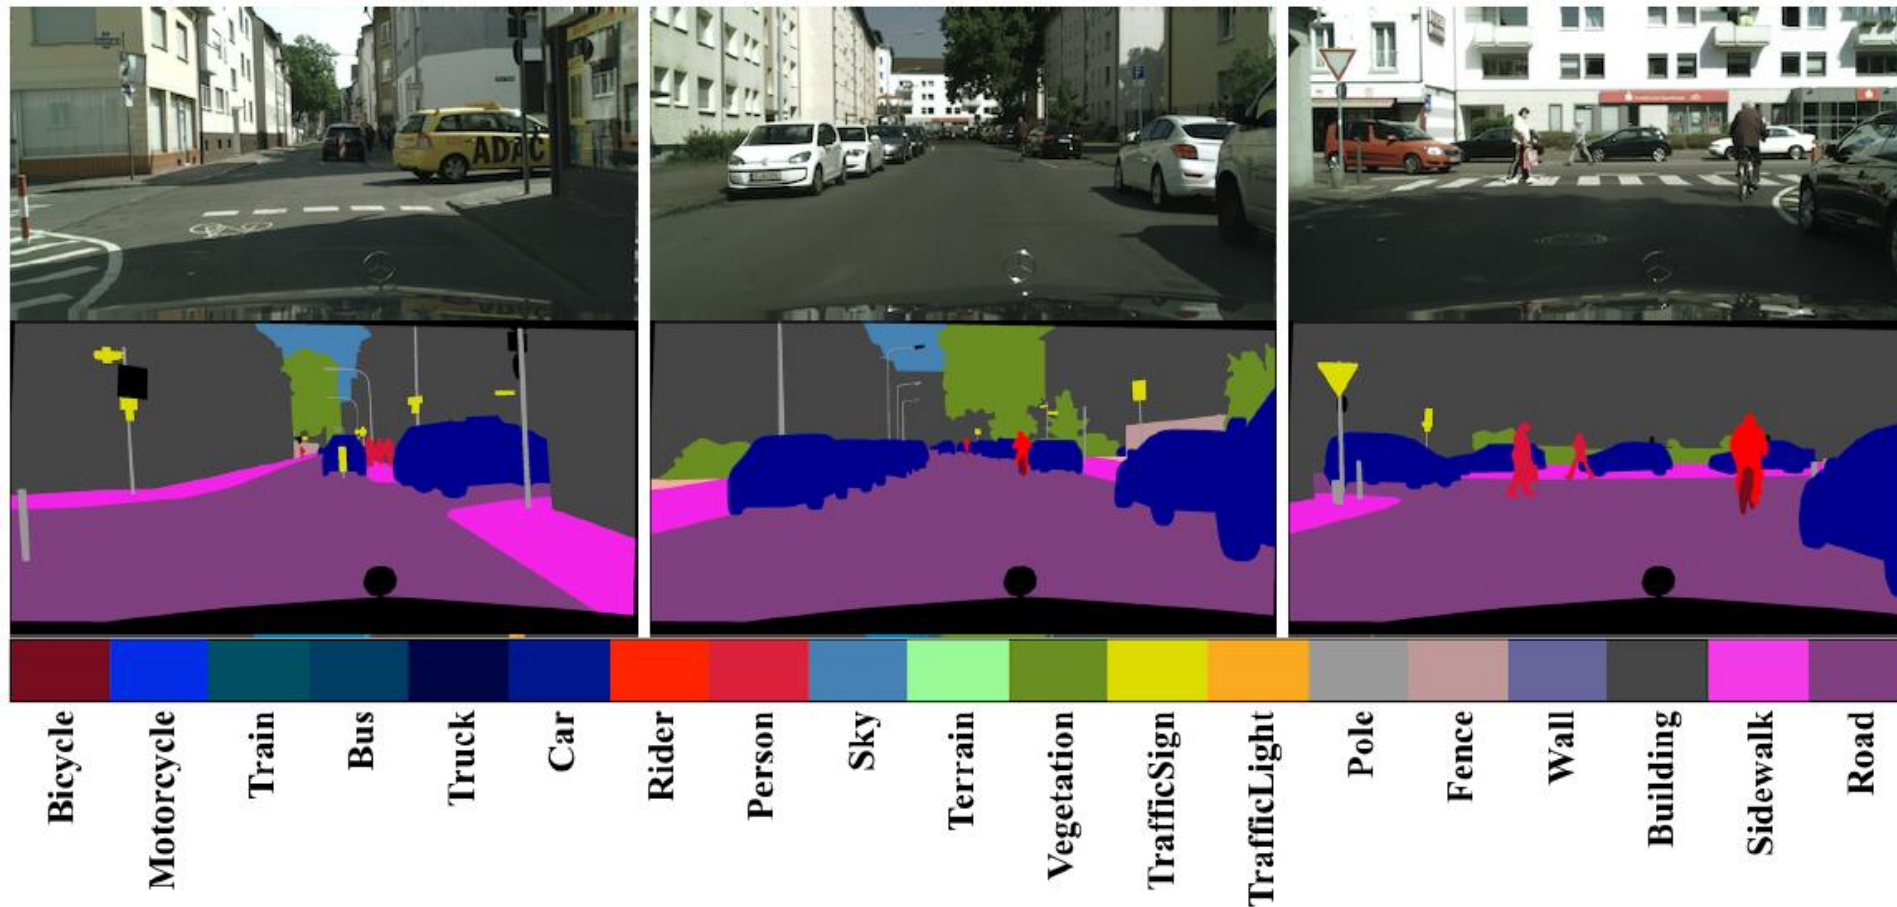
<https://emerj.com/ai-sector-overviews/how-self-driving-cars-work/>

Sensors Scheme and Functionality



<https://test.neurohive.io/en/state-of-the-art/self-driving-cars/>

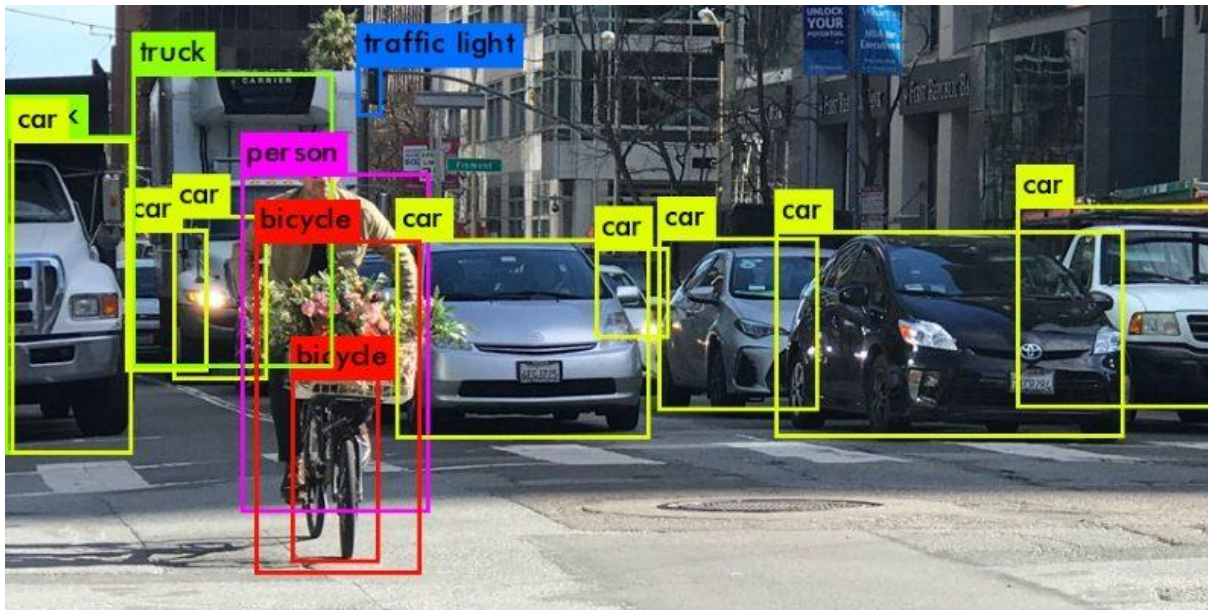
Role of Computer Vision and Image Processing



Taha Emara, [Hossam E. Abd El Munim](#), and Hazem M. Abbas, “LiteSeg: A Novel Lightweight ConvNet for Semantic Segmentation”, International Conference on Digital Image Computing: Techniques and Applications (DICTA), Perth, Australia, Dec, 2019

What is Computer Vision?

- Deals with the development of the theoretical and algorithmic basis by which useful information about the 3D world can be automatically extracted and analyzed from a single or multiple 2D images of the world.



<https://azati.ai/image-detection-recognition-and-classification-with-machine-learning/>



He et al, "Mask R-CNN", ICCV 2017

Computer Vision, Also Known As ...

- Image Analysis
- Scene Analysis
- Image Understanding

Some Related Fields are

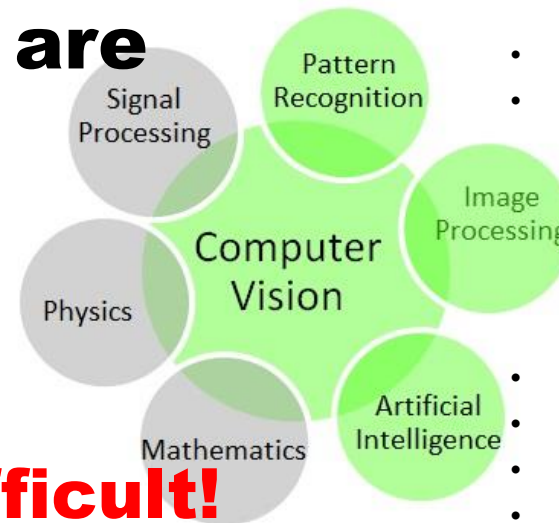
- Image Processing
- Computer Graphics
- Pattern Recognition
- Robotics
- Artificial Intelligence

Computer Vision is Difficult!

- Mapping 2D information to the real world is not unique.
- It is a computationally expensive process.
- Pose, illumination, scale variations always represent big challenges.

Applications

- Industrial inspection/quality control
- Surveillance and security
- Face recognition
- Gesture recognition
- Space applications
- Medical image analysis
- Autonomous vehicles
- Virtual reality and much more



Visual Recognition/Image Classification Challenges

Viewpoint variation



Scale variation



Deformation



Occlusion



Illumination conditions



Background clutter



Intra-class variation









<http://cs231n.github.io/classification/>

Top IEEE/Springer/IET Journals

Rank	Publisher	Journal Details
10	IEEE	IEEE Transactions on Pattern Analysis and Machine Intelligence ISSN:0162-8828 , Monthly
44	IEEE	IEEE Transactions on Image Processing ISSN:1057-7149 , Monthly
51	IEEE	IEEE Transactions on Information Forensics and Security ISSN:1556-6013 , Monthly
63	IEEE	IEEE Transactions on Robotics ISSN:1552-3098 , Bimonthly
68	IEEE	IEEE Transactions on Medical Imaging ISSN:0278-0062 , Monthly
83	IEEE	IEEE Transactions on Circuits and Systems for Video Technology ISSN:1051-8215 , Monthly
104	IEEE	IEEE Robotics and Automation Magazine ISSN:1070-9932 , Quarterly
139	IEEE	IEEE Journal of Selected Topics in Applied Earth Observations and Remote Sensing ISSN:1939-1404 , Monthly
156	IEEE	IEEE Geoscience and Remote Sensing Letters ISSN:1545-598X , Monthly
193	IEEE	IEEE Transactions on Audio, Speech and Language Processing ISSN:1558-7916 , Monthly
11	Springer	International Journal of Computer Vision ISSN:0920-5691 , Monthly
258	Springer	Journal of Real-Time Image Processing ISSN:1861-8200 , Quarterly
261	Springer	Machine Vision and Applications ISSN:0932-8092 , Bimonthly
265	Springer	Journal of Mathematical Imaging and Vision ISSN:0924-9907 , Monthly
269	Springer	Eurasip Journal on Advances in Signal Processing ISSN:1687-6180 , Irregular

Top Conferences

Hindex	Publisher	Conference Details
1	158 	CVPR : IEEE Conference on Computer Vision and Pattern Recognition, CVPR Jun 15, 2019 - Jun 21, 2019 - Long Beach , United States http://cvpr2019.thecvf.com/
5	89 	ICCV : IEEE International Conference on Computer Vision Oct 27, 2019 - Nov 3, 2019 - Seoul , South Korea http://iccv2019.thecvf.com/
36	50 	IROS : IEEE/RSJ International Conference on Intelligent Robots and Systems Nov 3, 2019 - Nov 8, 2019 - Macao , China http://www.iros2019.org/
94	34 	ICIP : IEEE International Conference on Image Processing Sep 22, 2019 - Sep 25, 2019 - Taipei , Taiwan http://2019.ieeeicip.org/
137	28 	FG : IEEE International Conference on Automatic Face & Gesture Recognition May 14, 2019 - May 18, 2019 - Lille , France http://fg2019.org/
146	27 	BTAS : IEEE International Conference on Biometrics: Theory Applications and Systems (BTAS) Sep 10, 2018 - Sep 13, 2018 - Los Angeles , United States https://www.isi.edu/events/btas2018/home
170	25 	ICME : IEEE International Conference on Multimedia and Expo Jul 23, 2018 - Jul 23, 2018 - San Diego , United States http://www.icme2018.org/important_dates

Hindex	Publisher	Conference Details
3	98 	ECCV : European Conference on Computer Vision Sep 8, 2018 - Sep 14, 2018 - Munich , Germany https://eccv2018.org/
55	43 	BMVC : British Machine Vision Conference Sep 5, 2018 - Sep 8, 2018 - Northumbria University , United Kingdom https://www.northumbria.ac.uk/about-us/news-events/events/2018/09/british-machine-vision-conference/
88	35 	ACCV : Asian Conference on Computer Vision Dec 2, 2018 - Dec 6, 2018 - Perth Western , Australia http://accv2018.net/
98	34 	MICCAI : Medical Image Computing and Computer Assisted Intervention Sep 16, 2018 - Sep 20, 2018 - Granada , Spain http://www.miccai2018.org/en/
319	15 	ISVC : International Symposium on Visual Computing Dec 12, 2016 - Dec 14, 2016 - Las Vegas , United States http://www.isvc.net
365	13 	ICIAP : International Conference on Image Analysis and Processing Sep 11, 2017 - Sep 15, 2017 - Catania , Italy http://www.iciap2017.com/

Machine Learning Problems

Supervised Learning

Unsupervised Learning

Discrete

classification or
categorization

clustering

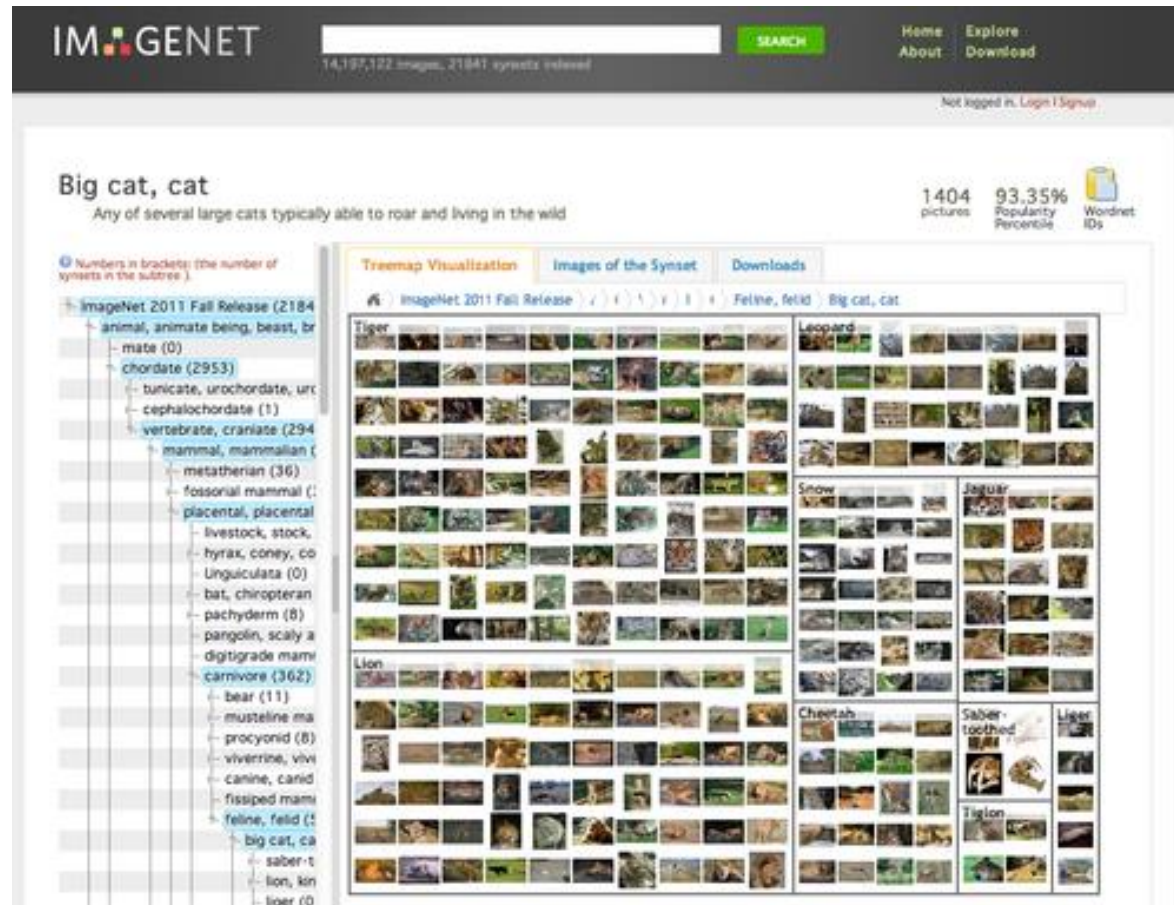
Continuous

regression

dimensionality
reduction

ImageNet

- Images for each category of WordNet
- 1000 classes
- 1.2mil images
- 100k test



The machine learning framework

- Apply a prediction function to a feature representation of the image to get the desired output:

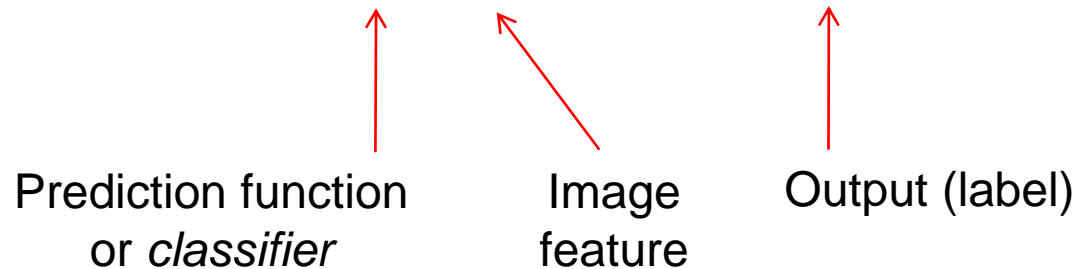
$f(\text{apple image}) = \text{"apple"}$

$f(\text{tomato image}) = \text{"tomato"}$

$f(\text{cow image}) = \text{"cow"}$

The machine learning framework

$$f(\mathbf{x}) = y$$



Training: Given a *training set* of labeled examples:

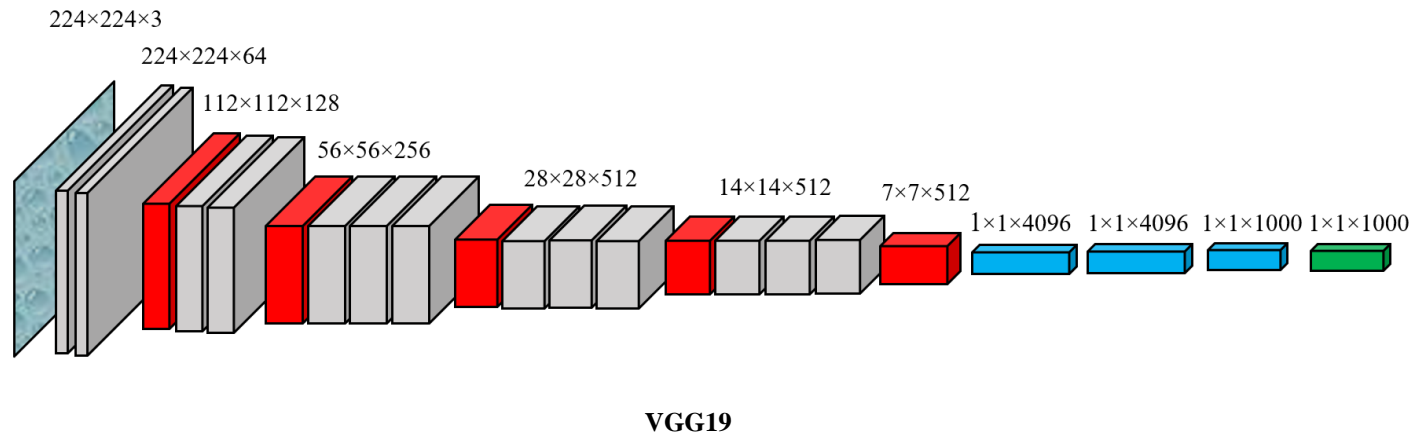
$$\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)\}$$

Estimate the prediction function f by minimizing the prediction error on the training set.

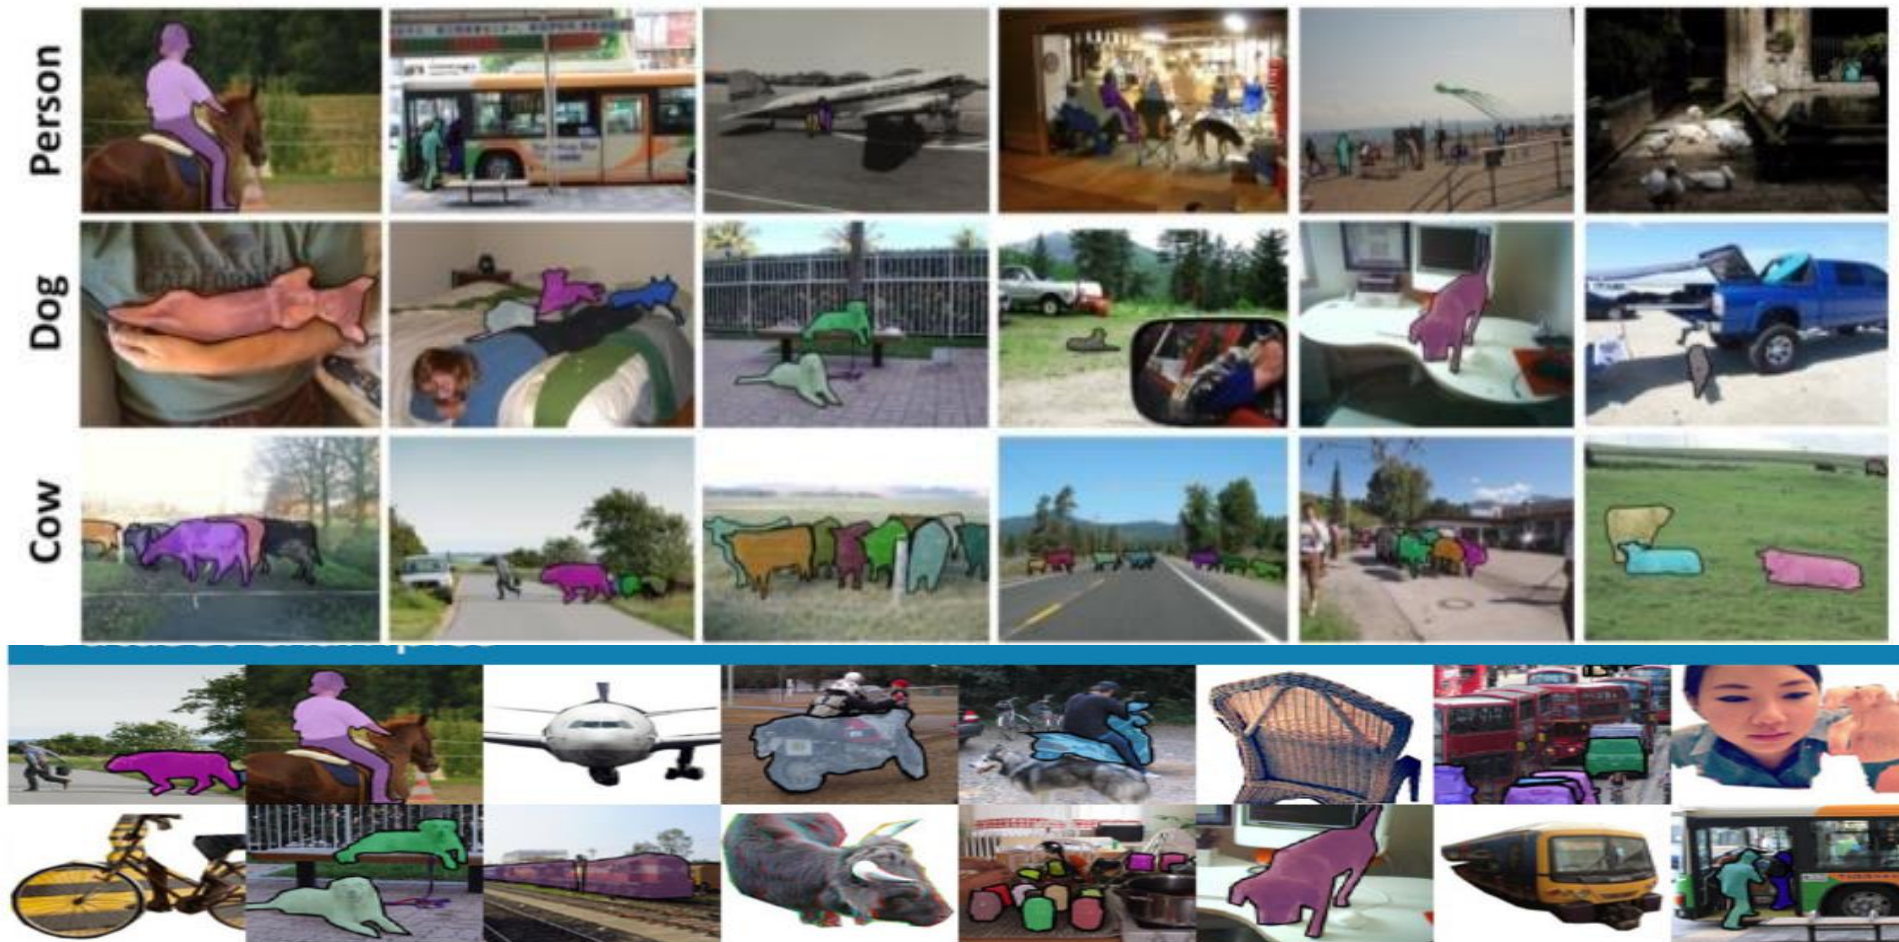
Testing: Apply f to a unseen *test example* \mathbf{x}_u and output the predicted value $y_u = f(\mathbf{x}_u)$ to *classify* \mathbf{x}_u .

State of the Art

- With enough training data, computer vision nearly matches human vision at most recognition tasks
- Deep learning has been an enormous disruption to the field. Many techniques are being “deepified”.
- The world of computer vision changed when deep learning arrived. For most of the computer vision tasks, deep learning models (Convolutional Neural Network) were built and trained which started outpacing the counter-part old machine learning methods.
- Applications are: Image classification, Object detection, Object tracking, Pose estimation, Text detection & recognition, Visual saliency detection, Action recognition, Scene labelling



Microsoft COCO: Common Objects in Context



Object Detection



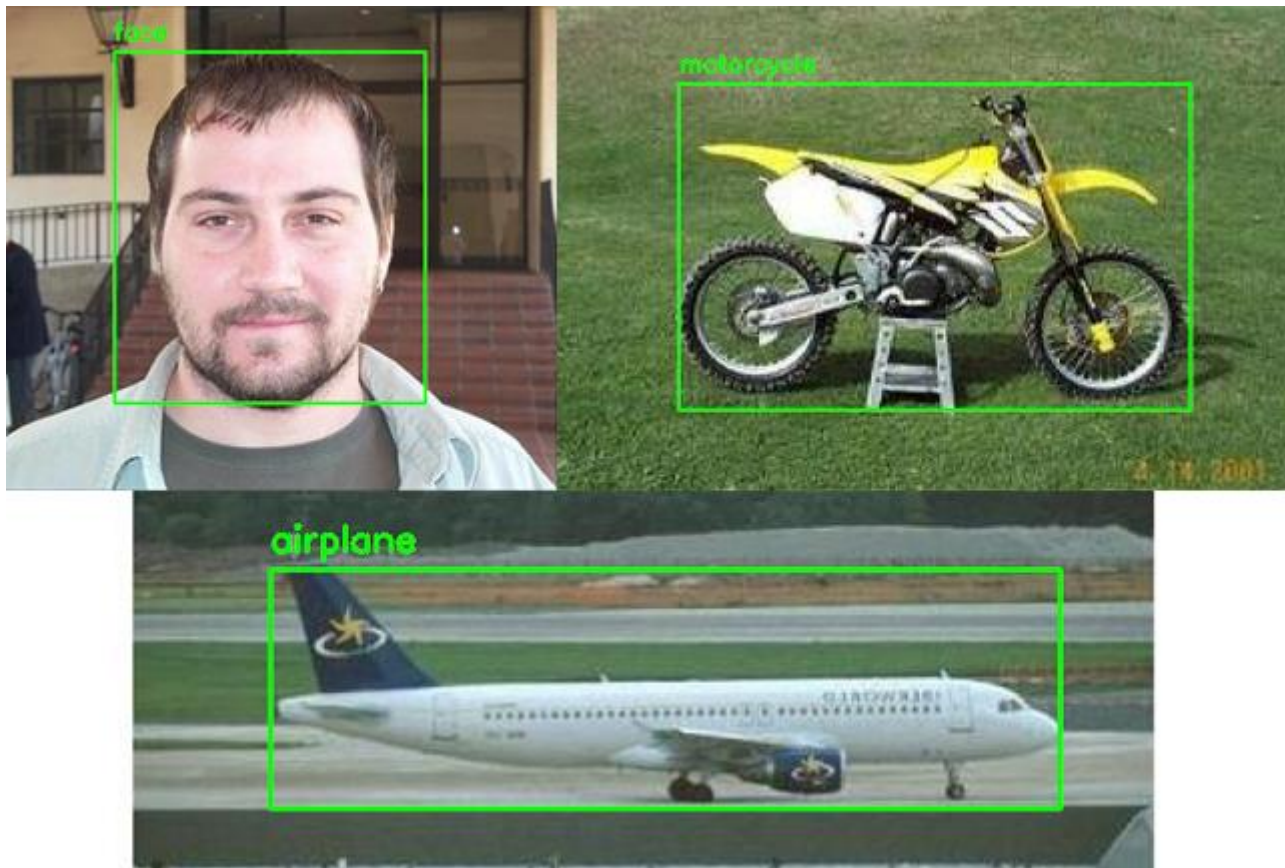
The image shows a busy city street scene with various objects detected and labeled with bounding boxes and text. The labels include:

- traffic light
- car
- person
- handbag
- bicycle
- backpack
- truck

The scene is a busy city street with tall buildings, traffic lights, and pedestrians. The labels are placed over the corresponding objects in the image.

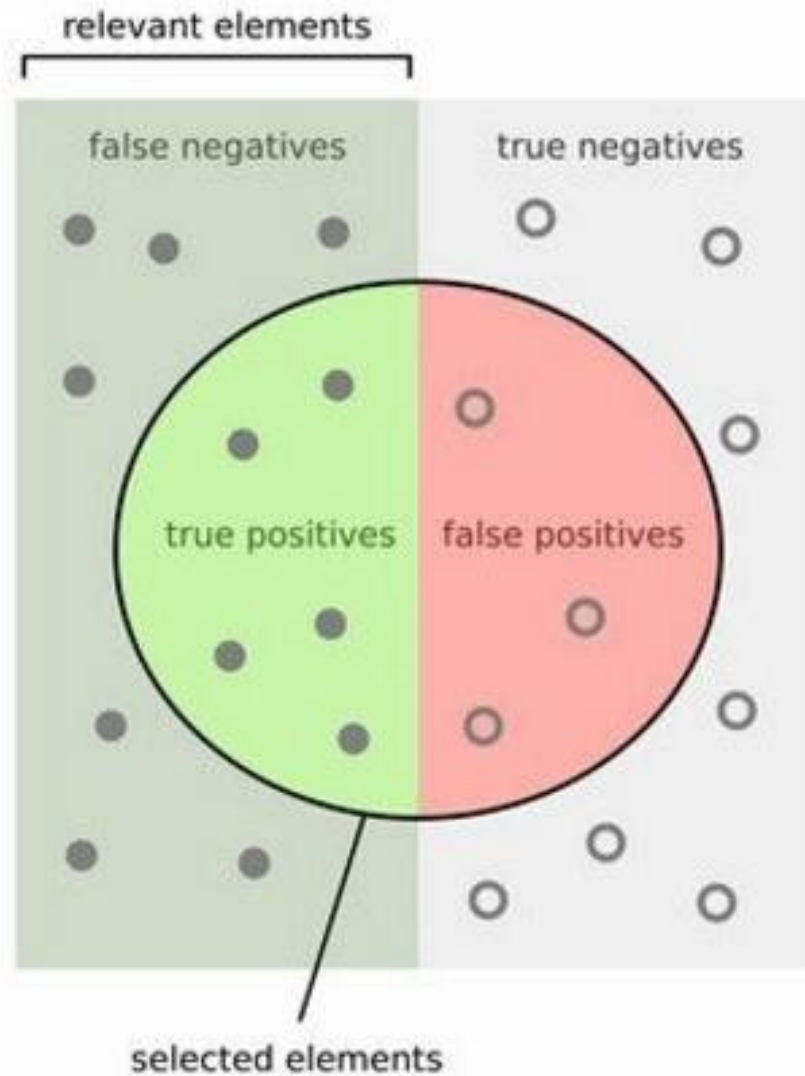


Regression/Classification and Object Detection



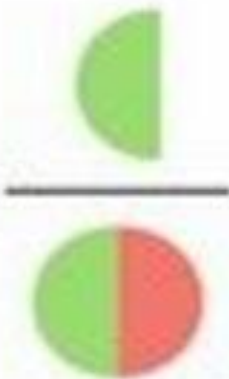
It aims to find the box parameters which are considered analog.

TP/FP

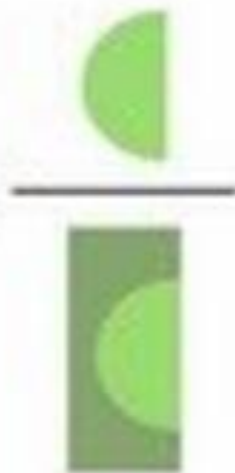


Precision

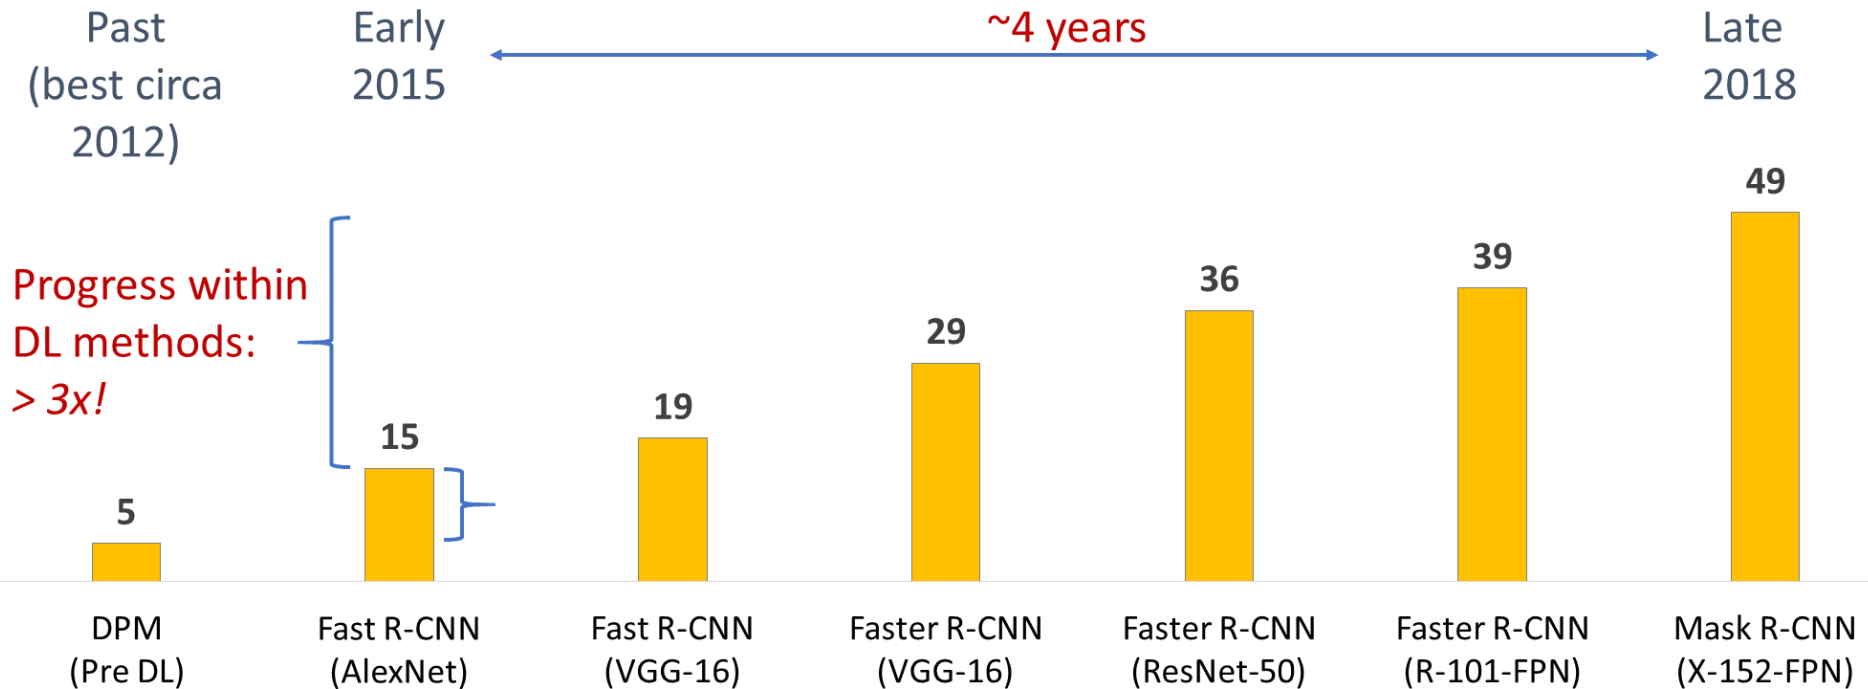
How many selected items are relevant?

$$\text{Precision} = \frac{\text{Relevant}}{\text{Selected}}$$


How many relevant items are selected?

$$\text{Recall} = \frac{\text{Relevant}}{\text{Total Relevant}}$$


COCO Object Detection Average Precision (%)



Dataset examples



Microsoft COCO: Common Objects in Context

25 October 2020

Hossam Abdelmunim

23

Why These Improvements in Performance?

- Features are learned rather than hand-crafted.
- More layers capture more invariances¹.
- More data to train deeper networks.
- More computing (GPUs).
- Better regularization: Dropout.
- New nonlinearities (Max pooling, Rectified linear units (ReLU)).
- Better optimization techniques².
- Several learning libraries have emerged as winners which provide a lot of support and convenience to train deep learning models for visual recognition and other visual tasks

theano

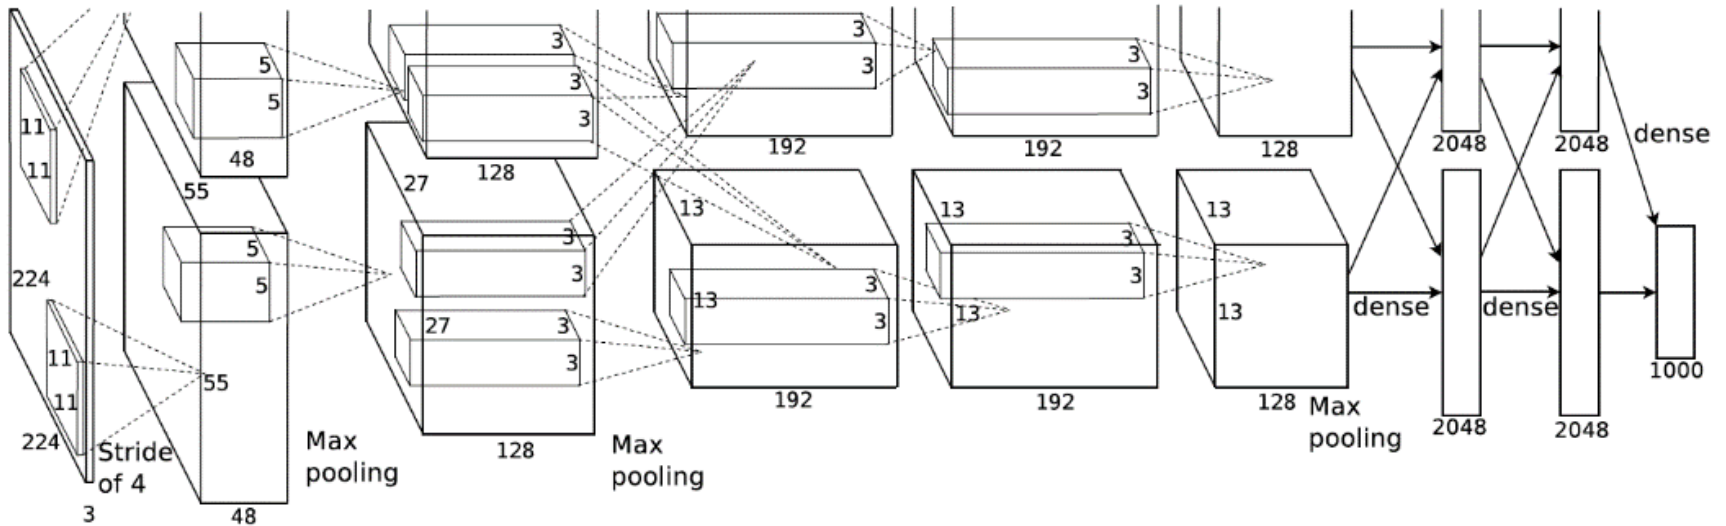


[1] Razavian, Azizpour, Sullivan, Carlsson, CNN Features off-the-shelf: an Astounding Baseline for Recognition. CVPRW'14.

[2] Diederik P. Kingma and Jimmy Lei Ba. Adam: a Method for Stochastic Optimization. International Conference on Learning Representations, pages 1–13, 2015.

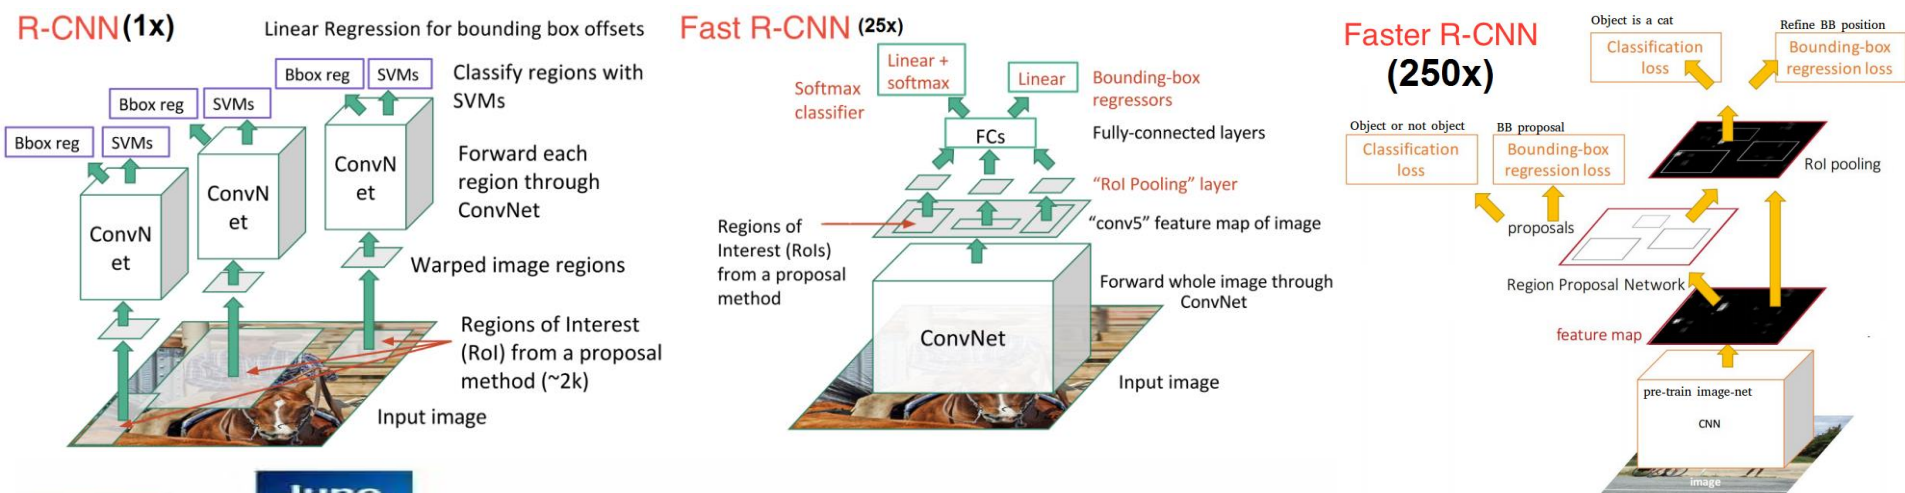
The way you see the world will be changed by the following CNN-based vision applications.

1- Image Classification



Since AlexNet (The winner of the 1st ImageNet competition, Alex Krizhevsky (Neural Information Processing - NIPS 2012)), there have been multiple new models using CNN as their backbone architecture and achieving excellent results in ImageNet: [ZFNet](#) (2013), [GoogLeNet](#) (2014), [VGGNet](#) (2014), Residual Net [ResNet](#) (2015), [DenseNet](#) (2016) etc.

2- Object Detection (CVPR-2014, ICCV-2015, PAMI-2017)

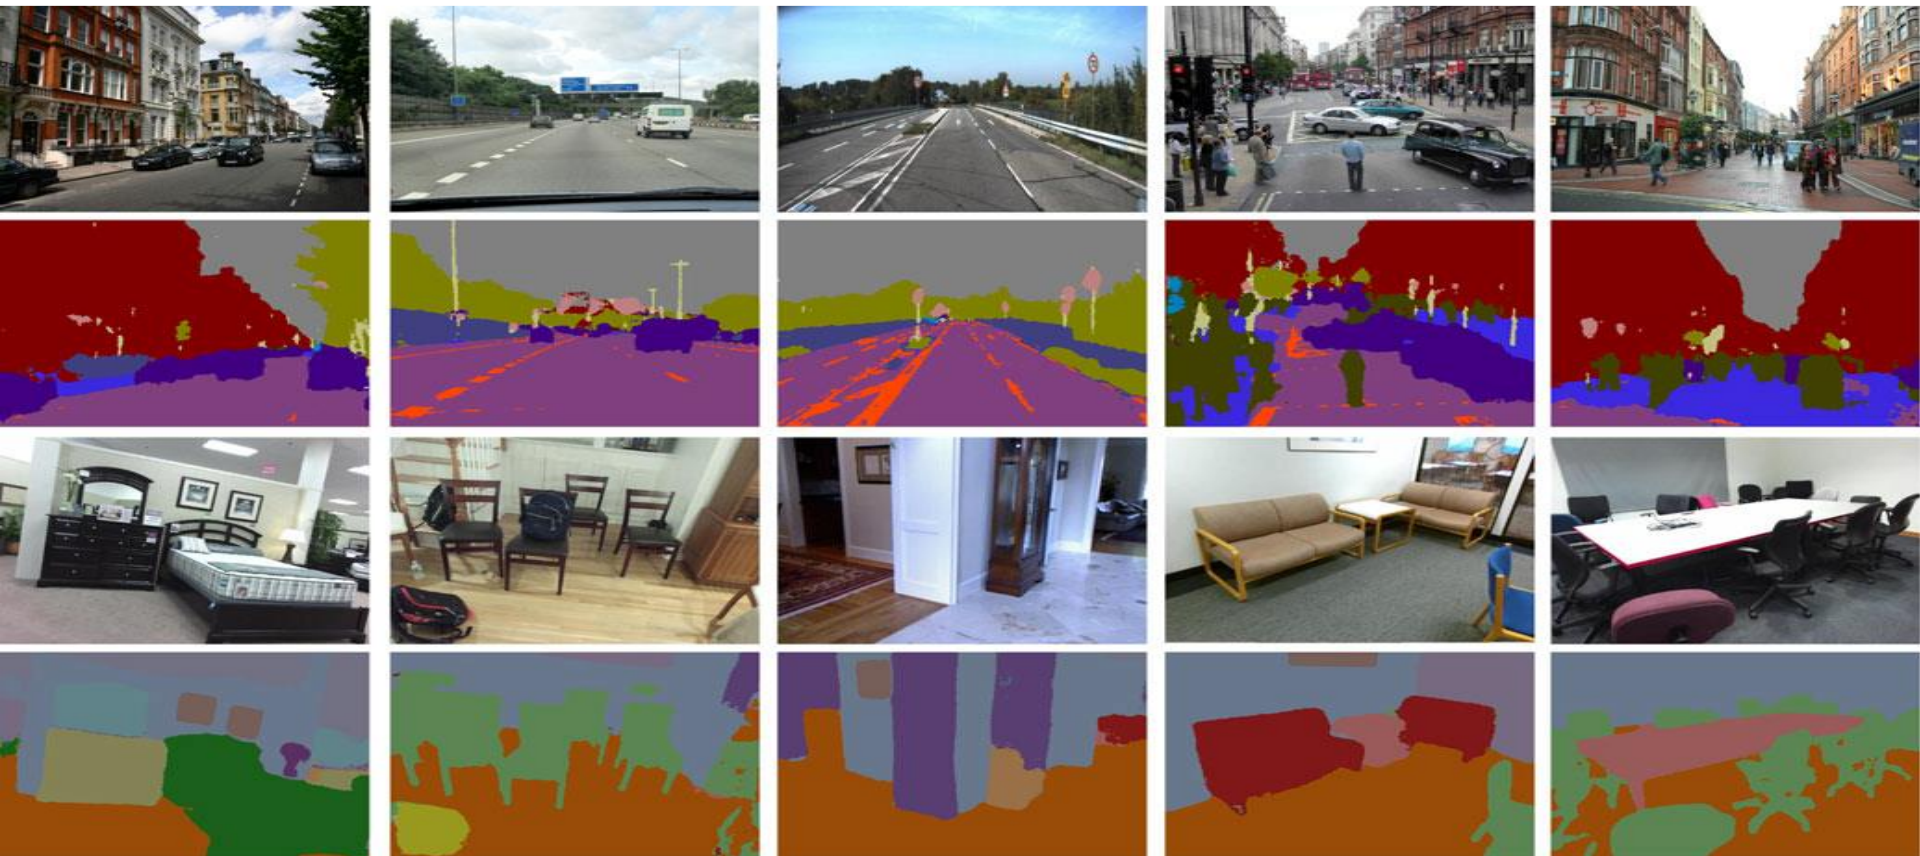
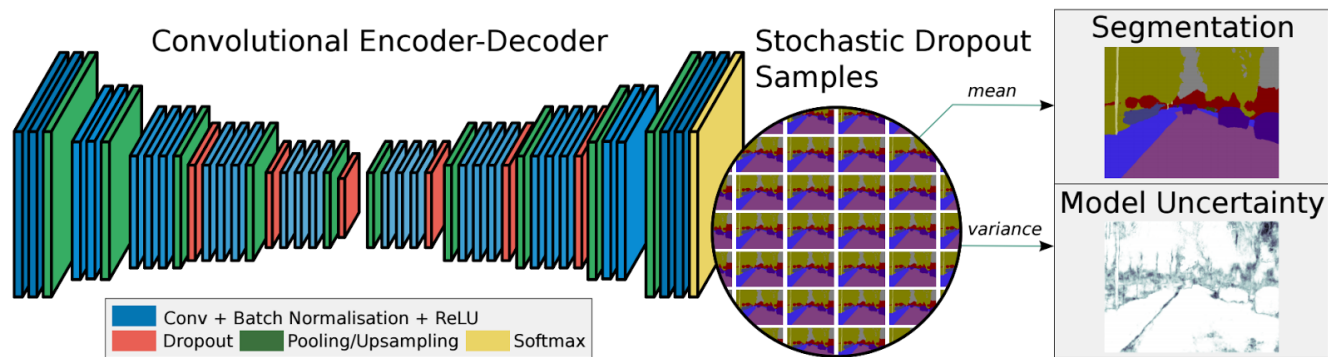


3- Object Tracking

Learning Multi-Domain Convolutional Neural Networks for Visual Tracking (CVPR 2016)



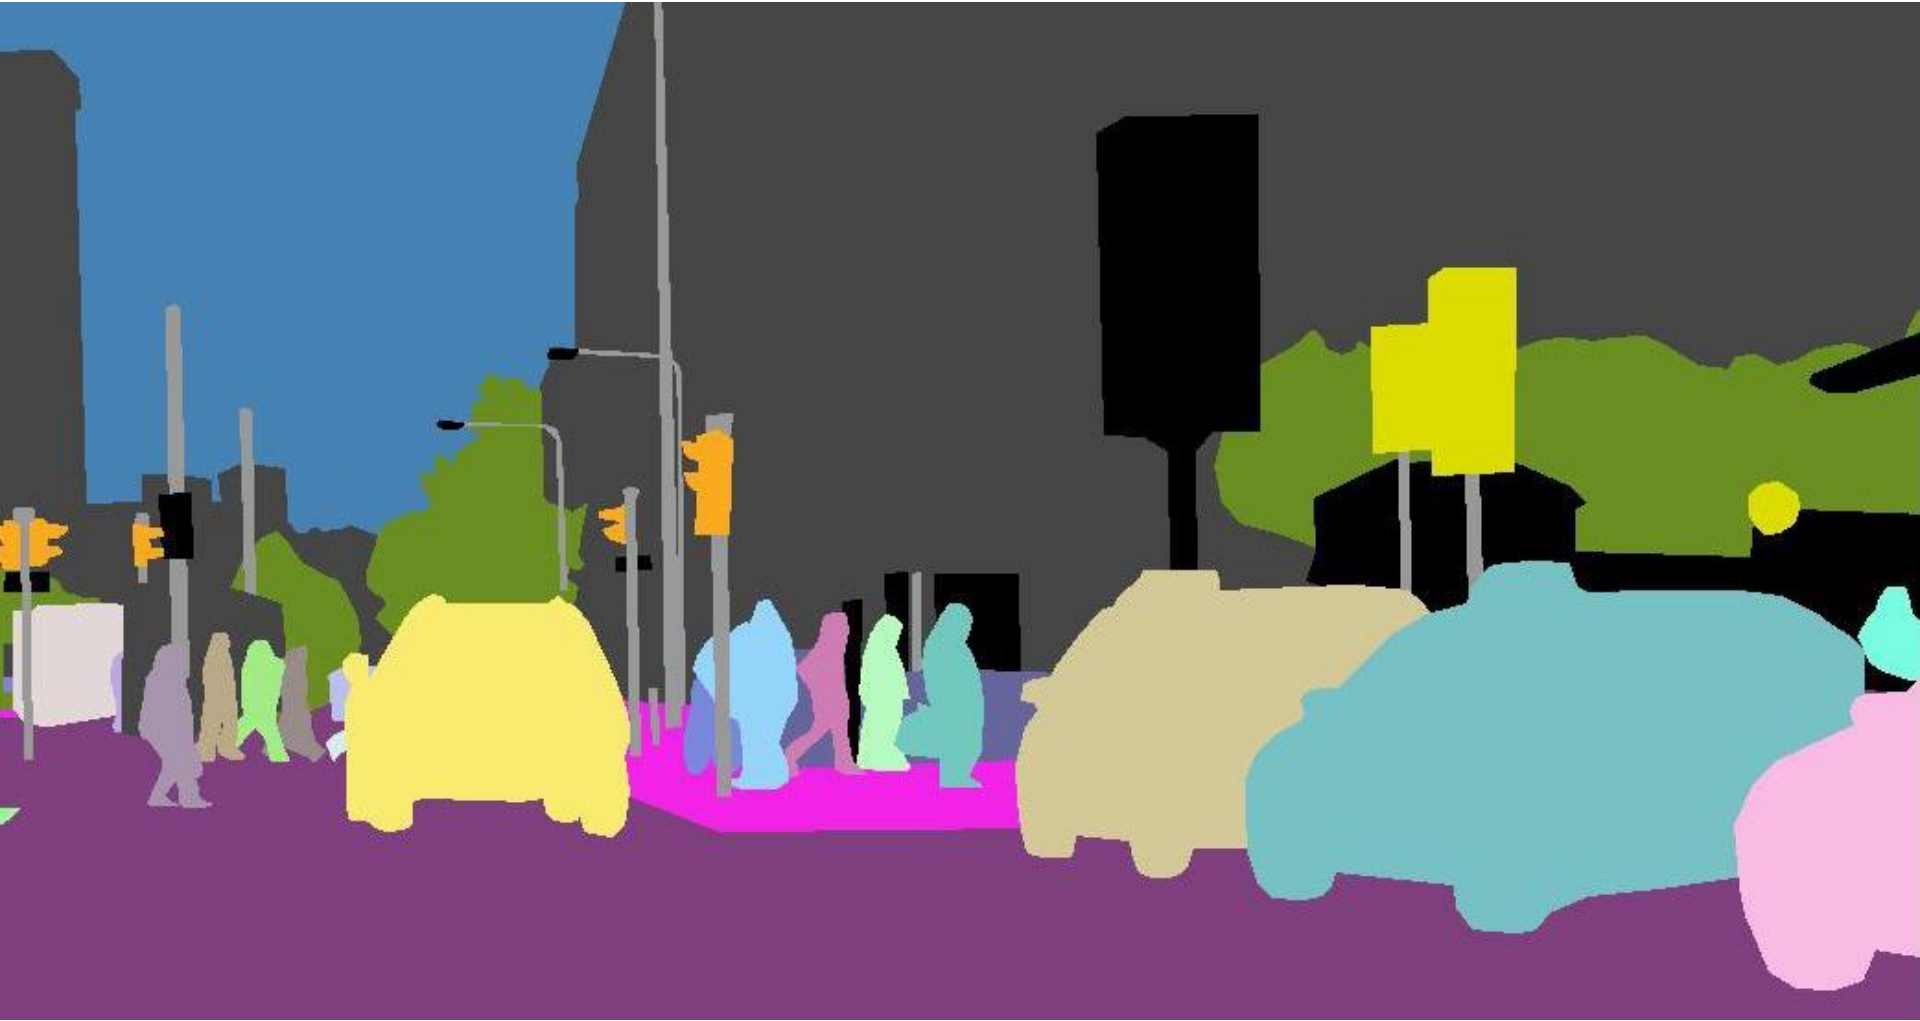
4- Semantic Segmentation: SegNet - IEEE PAMI 2017



Recent research in Semantic Segmentation all relies heavily on fully convolutional networks, such as [Dilated Convolutions](#), [DeepLab](#) (IEEE PAMI 2018), and [RefineNet](#).

5- Instance Segmentation:

Mask R-CNN (arXiv:1703.06870v3 [cs.CV] 24 Jan 2018)



Some Non-CNN Based Techniques Still Evolve

- IEEE PAMI 2019, Image Projective Invariants
- IEEE PAMI 2018, Towards Reaching Human Performance in Pedestrian Detection
- IEEE PAMI 2018, SIFT Meets CNN: A Decade Survey of Instance Retrieval
- ICCV 2017, Reconfiguring the Imaging Pipeline for Computer Vision
- CVPR 2018: Five-point Fundamental Matrix Estimation for Uncalibrated Cameras
- CVPR 2018: Single View Stereo Matching